




上海交通大学  
SHANGHAI JIAO TONG UNIVERSITY



# 游戏设计与开发

GPU实时图形管线  
GPU: Real-time Graphics Pipelines


上海交通大学软件学院  
数字艺术实验室  
Digital ART LAB



## Outline


- Real Time Requirement
- A Conceptual Rendering Pipeline
- Evolution of GPU
  - Hardware
  - APIs
  - Shaders

Digital ART Lab, SJTU



## REAL TIME REQUIREMENT

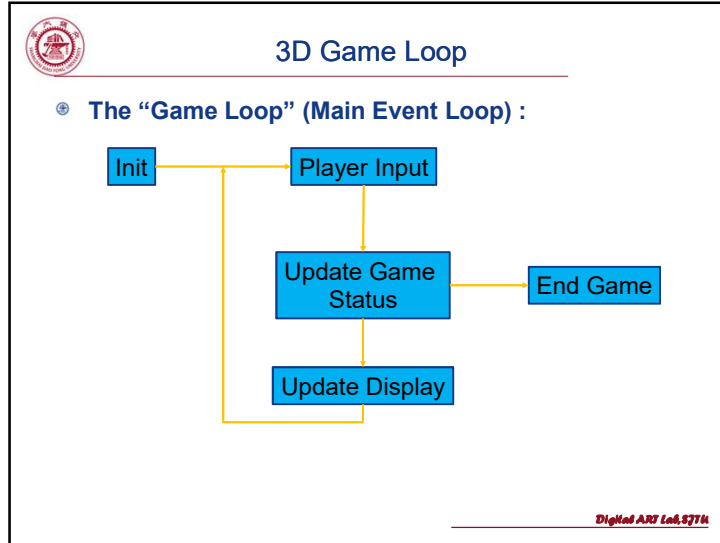
Digital ART Lab, SJTU



## Building A Game

- Background Scene (e.g. sky, terrain)
- Static Objects
- Movement of Objects
- Users' Control
- Collision and Response
- Others

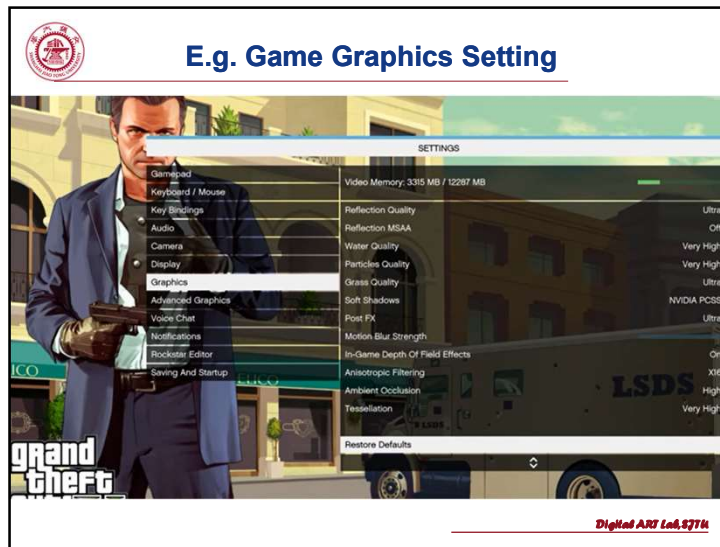
Digital ART Lab, SJTU



### How Fast Does my Game Loop Need to Run?

- ANSWER: “*It depends...*”
  - Visual displays: 25-30 Hz or higher, 90~120Hz for VR display
  - Head-tracking for HMDs: 60 Hz, but even only 2-5ms of latency yields *display lag*, which often quickly causes users to lose their lunch...
  - Haptic displays: much higher update rates (500 - 1000 Hz)
  - Multitasking / Multiprocessing: allows for different update rates for different types of output displays
- Main requirement: **Real Time 3D Graphics**

Digital ART Lab, SJTU





## VR in VR (IEEE VR 2020, Mar.23-26)



Digital ART Lab, SJTU



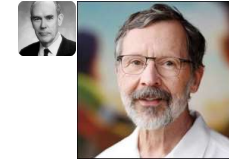
## ACM A.M. Turing Award (2019)

### PIONEERS OF MODERN COMPUTER GRAPHICS: 3D CGIs



Patrick M. Hanrahan

- Volume rendering; light field rendering
- RenderMan Shading Language: Shader, → GLSL
- Brook (a language for GPUs) → Nvidia's CUDA
- GPU enabled:
  - CGI for animated films and games
  - VR/AR
  - high performance computing
  - machine learning for AI



Edwin E. Catmull

Curved patches displaying  
Z-Buffering  
Texture Mapping  
Catmull-Clark Subdivision

Lucas Films → Pixar (1986)

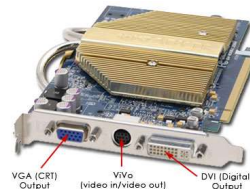
Digital ART Lab, SJTU



## Graphics Hardware and APIs

### GPU(Graphics Processing Unit):

- Nvidia, AMD, Intel



### API:

- Microsoft: DirectX 9, DirectX 10, DirectX 11,...
- OpenGL ARB: OpenGL 1.0, 2.0, 3.0, 4.0,...

Digital ART Lab, SJTU



## 3D Graphics Software Tools

- Low-level 3D graphics APIs:
  - hardware-independent, transparent access
    - window-system and OS independent
    - network-transparent
  - Commonly-Used 3D Graphics APIs:
    - Realtime: OpenGL, Direct3D (part of DirectX), Vulkan
    - Offline: Renderman
- High-level 3D graphics tools:
  - 3D Graphics Engine: OGRE, OpenSceneGraph, VTK
  - 3D Game Engine: Unity, Unreal

Digital ART Lab, SJTU



## OpenGL vs. Vulkan

### Vulkan: lower-level; multi-core; parallel tasking

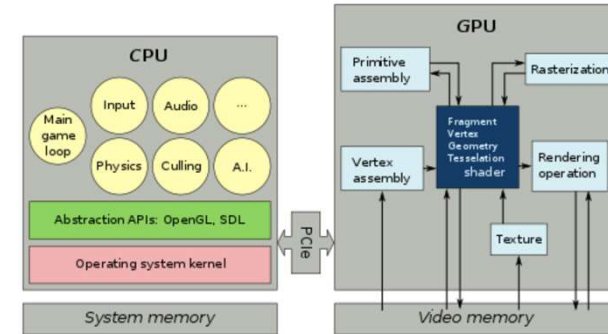
- lower overhead
- more direct control over the GPU
- and lower CPU usage

OpenGL	Vulkan
One single global state machine	Object-based with no global state
State is tied to a single context	All state concepts are localized to a command buffer
Operations can only be executed sequentially	Multi-threaded programming is possible
GPU memory and synchronization are usually hidden	Explicit control over memory management and synchronization
Extensive error checking	Vulkan drivers do no error checking at runtime; there is a validation layer for developers

Digital ART Lab, SJTU



## CPU-GPU Cooperation



Digital ART Lab, SJTU



## GPU vs. CPU

### GPU

- Data Parallelism
- More cores, but low clock speed
- VRAM, fast but small size
- Low cache memory



### CPU

- Task parallelism
- Less cores, but high clock speed
- External RAM, slow but large in size
- High cache memory



Digital ART Lab, SJTU



## Graphics Computations on CPU

- Vertices to pixels:
    - Transformations done on CPU
    - Compute each pixel in series...slow!
  - Example:
    - 1 million triangles
    - x 100 pixels per triangle
    - x 10 lights
    - x 4 cycles per light computation
- = 4 billion cycles

Digital ART Lab, SJTU



## Flynn's Taxonomy

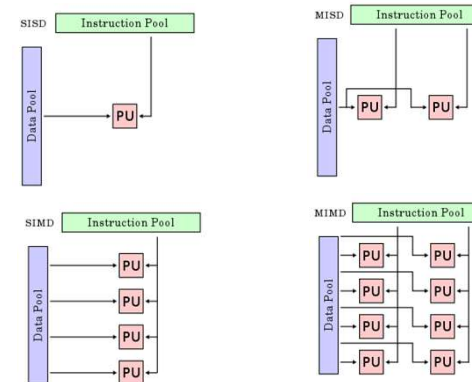
- ⊕ A taxonomy of computer architecture
  - Micheal Flynn in 1966
- ⊕ Based on two things:
  - Instructions
  - Data

	Single instruction	Multiple instruction
Single data	<b>SISD</b>	<b>MISD</b>
Multiple data	<b>SIMD</b>	<b>MIMD</b>

Digital ART Lab, SJTU



## Which one is closest to GPU?



Digital ART Lab, SJTU



## A CONCEPTUAL RENDERING PIPELINE

Digital ART Lab, SJTU



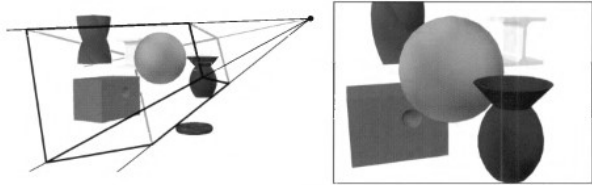
## The Graphics Rendering Pipeline

- rendering pipeline or the pipeline
  - The core component of real time graphics
  - main function: generate, or render, a 2D image, given a virtual camera, 3D objects, light sources, lighting models, textures, and more.

Digital ART Lab, SJTU

### Using the pipeline

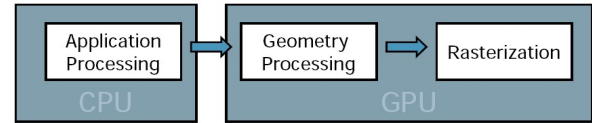
- Object (model, color, texture, material)
- Camera/view
- Light source



*Digital ART Lab, SJTU*

### The Basic Construction

- A conceptual rendering pipeline
- 3 conceptual stages:



*Digital ART Lab, SJTU*

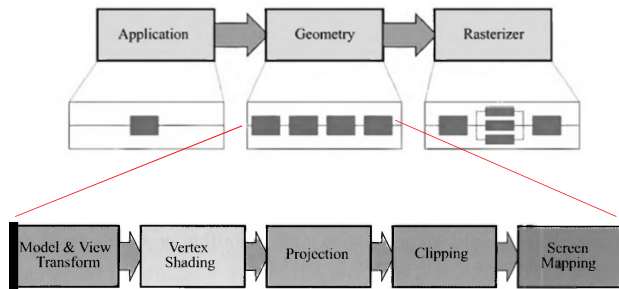
### High-Level Pipeline

Application	Geometry (a.k.a. "vertex pipeline")	Rasterization (a.k.a. "pixel pipeline" or "fragment pipeline")
Handle input	Transform	Rasterize (fill pixels)
Simulation & AI	Lighting	Interpolate vertex parameters
Culling	Skinning	Look up/filter textures
LOD selection	Calculate texture coords	Z- and stencil tests
Prefetching		Blending

*Digital ART Lab, SJTU*

### Substages of Geometry Stage

- ⊕ A pipeline consists of several stages (N)  
— ideally give a speedup of a factor of N



*Digital ART Lab, SJTU*

### Model Transform

■ model transform: model space  $\rightarrow$  world space

Model & View Transform  $\rightarrow$  Vertex Shading  $\rightarrow$  Projection  $\rightarrow$  Clipping  $\rightarrow$  Screen Mapping

*Digital ART Lab, SJTU*

### View Transform

- ⊕ Purpose: place camera at the origin and aim it to look in  $-z$ , with  $y$  upward,  $x$  to the right
- ⊕ Camera space (or eye space)
- ⊕ right-hand or left-hand (API specific)

Model & View Transform  $\rightarrow$  Vertex Shading  $\rightarrow$  Projection  $\rightarrow$  Clipping  $\rightarrow$  Screen Mapping

*Digital ART Lab, SJTU*

### Transform Matrix

⊕ Both are implemented as 4x4 matrices

— Model Transform

$$\begin{bmatrix} x' \\ y' \\ z' \\ 1 \end{bmatrix} = \begin{bmatrix} 1 & 0 & 0 & t_x \\ 0 & 1 & 0 & t_y \\ 0 & 0 & 1 & t_z \\ 0 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

— View Transform

$$\begin{bmatrix} x' \\ y' \\ z' \\ 1 \end{bmatrix} = \begin{bmatrix} a & b & c & d \\ e & f & g & h \\ i & j & k & m \\ 0 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

Model & View Transform  $\rightarrow$  Vertex Shading  $\rightarrow$  Projection  $\rightarrow$  Clipping  $\rightarrow$  Screen Mapping

*Digital ART Lab, SJTU*

### Vertex Shading

- ⊕ Modeling object appearance:
  - Object's materials
  - Effects of light sources
- ⊕ Shading :
  - Determining the effect of a light on a material
  - By computing a shading equation at points on the object
    - Per-vertex operation
    - or Per-pixel operation

Model & View Transform  $\rightarrow$  Vertex Shading  $\rightarrow$  Projection  $\rightarrow$  Clipping  $\rightarrow$  Screen Mapping

*Digital ART Lab, SJTU*





## Vertex Shading

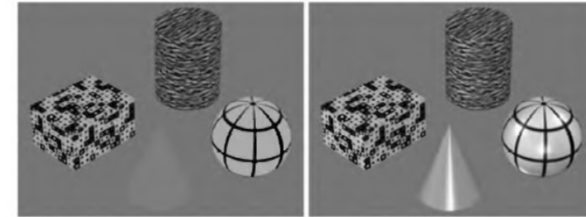
- ⊙ Material data can be stored at each vertex:
  - Points' location
  - A normal
  - A color
  - Other numerical information
- ⊙ Shading can be performed in
  - World space
  - Model space
  - Eye space

Digital ART Lab, SJTU



## Lighting

### ■ Lighting in fixed-function pipeline

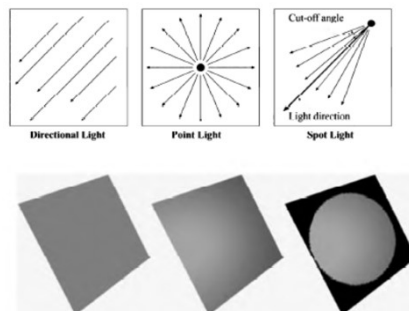


Digital ART Lab, SJTU



## Lighting

### ■ Light sources



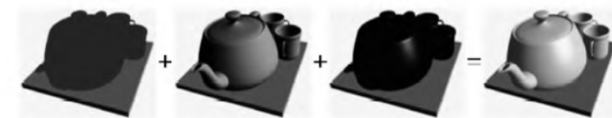
Digital ART Lab, SJTU



## Lighting

### ■ Lighting Equation in fixed-function pipeline

$$i_{tot} = i_{amb} + i_{diff} + i_{spec} \quad (4.14)$$



Digital ART Lab, SJTU



### Shading Model

- **Shading Model:**
  - Flat shading, Gouraud shading, Phong Shading (not in fixed)

triangle flat shading shading interpolation Gouraud shading Phong shading

Digital ART Lab, SJTU

### Projection

- View volume  $\rightarrow$  unit cube  $(-1,-1,-1)$  and  $(1,1,1)$
- Orthographic projection, perspective projection

Model & View Transform Vertex Shading Projection Clipping Screen Mapping

Digital ART Lab, SJTU

### Orthographic Projection

$$M_{\text{orth}} = \begin{bmatrix} \frac{2}{r-l} & 0 & 0 & -\frac{r+l}{r-l} \\ 0 & \frac{2}{t-b} & 0 & -\frac{t+b}{t-b} \\ 0 & 0 & \frac{2}{n-f} & -\frac{n+f}{n-f} \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

Digital ART Lab, SJTU

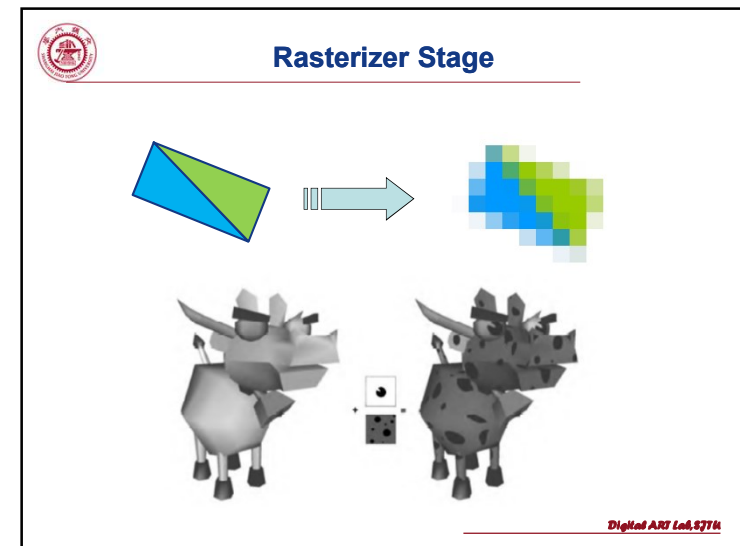
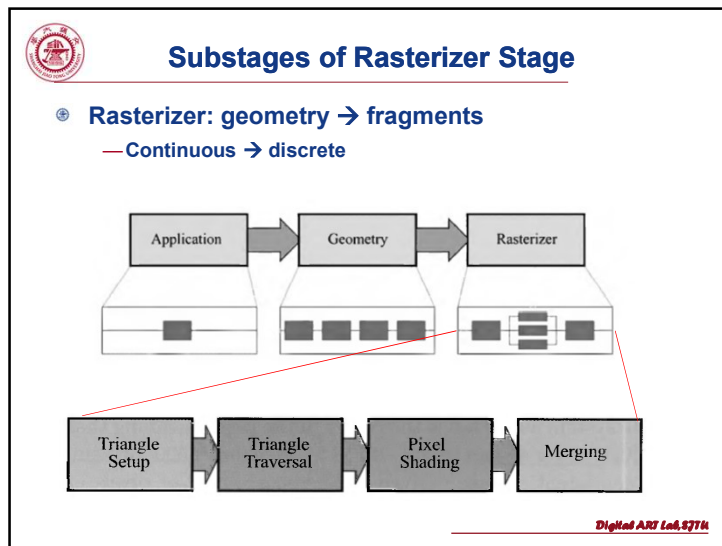
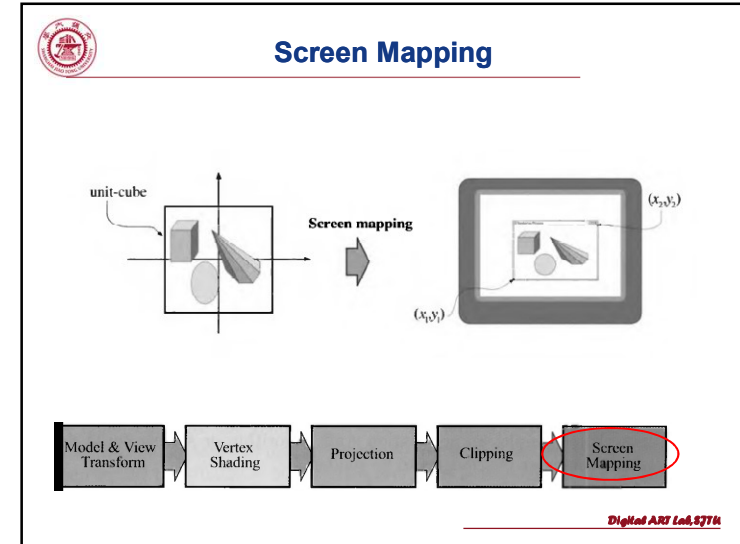
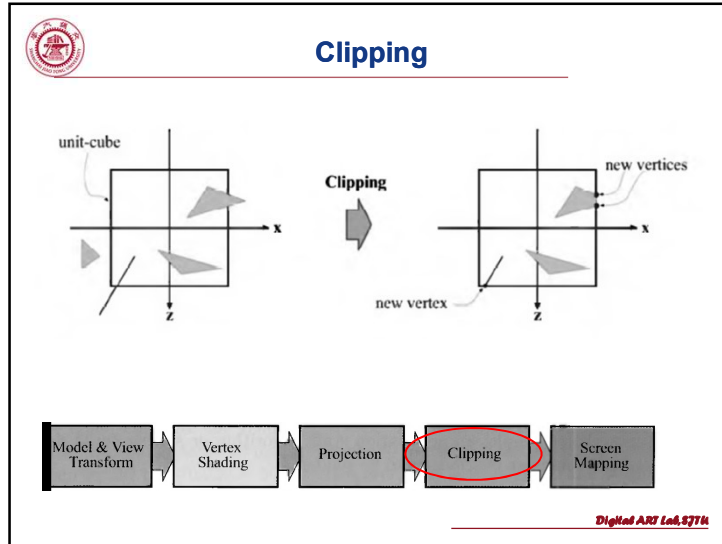
### Orthographic $\rightarrow$ Perspective

$$M_{\text{per}} = M_{\text{orth}} P$$

$$M_{\text{orth}} = \begin{bmatrix} \frac{2}{r-l} & 0 & 0 & -\frac{r+l}{r-l} \\ 0 & \frac{2}{t-b} & 0 & -\frac{t+b}{t-b} \\ 0 & 0 & \frac{2}{n-f} & -\frac{n+f}{n-f} \\ 0 & 0 & 0 & 1 \end{bmatrix} \quad M_{\text{per}} = \begin{bmatrix} \frac{2n}{r-l} & 0 & \frac{l+r}{l-r} & 0 \\ 0 & \frac{2n}{t-b} & \frac{b+t}{b-t} & 0 \\ 0 & 0 & \frac{f+n}{n-f} & \frac{2fn}{f-n} \\ 0 & 0 & 1 & 0 \end{bmatrix}$$

from Fundamentals of Computer Graphics, 4th Edition

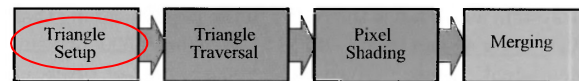
Digital ART Lab, SJTU





## Triangle Setup

- Compute the differentials and other data for the triangle's surfaces
  - used for scan conversion (next stage)
  - for interpolation of various shading data produced by the geometry stage (next stages)
- Performed by fixed-operation hardware

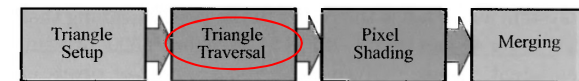


Digital ART Lab, SJTU



## Triangle Traversal

- Also called **scan conversion**
- Find which samples or pixels are inside a triangle
  - Check each pixel covered by the triangle
  - Generate a **fragment** for the part of the pixel that overlaps the triangle
  - Fragment properties: depth, and any shading data from the geometry stage:
    - generated using data interpolated among the 3 triangle vertices
- Performed by fixed-operation hardware

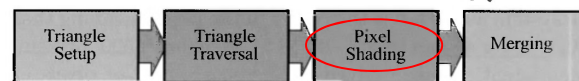


Digital ART Lab, SJTU



## Pixel Shading

- Do any **per-pixel shading** computations
  - Input: the interpolated shading data
  - Output: one or more colors to be passed to next stage
- Executed by programmable GPU cores
- Many techniques can be employed here, e.g.
  - Texturing: “glue” an image (1D, 2D or 3D) onto the object

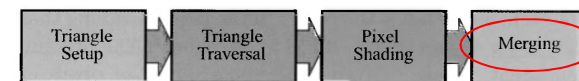


Digital ART Lab, SJTU



## Merging

- Combine the fragment color produced with the color currently stored in color buffer
  - **color buffer**: 2D array store color info for each pixel (e.g. RGB)
- Not fully programmable, but highly configurable (enable SFXs)
- Also resolving visibility: mostly **depth test** using **Z-buffer** (same size as the color buffer, store depth value)
  - Depth test: compare z-value before rendering to a pixel → update Z-buffer and color buffer if closer
  - **Order-independent** for opaque object, but need back-to-front for transparent object (major weakness of Z-buffer)



Digital ART Lab, SJTU



## Merge stage: more buffers and operations

- ④ Color buffer: colors
- ④ Z-buffer: z-values (*depth test*)
- ④ Alpha channel (color buffer): opacity values
  - *alpha test* ( $=$ ,  $>$ , ...) optional before the *depth test*
  - E.g. ensure fully transparent fragments not affect z-buffer
- ④ Stencil buffer
  - Record locations of the rendered primitive
  - offscreen buffer (typically 8 bits/pixel)
  - Control rendering into the color buffer and Z-buffer
  - Powerful tool for SFXs: e.g. a circle window
  - Raster operations (ROP) or blend operations
- ④ Frame buffer (all the buffers, or color + Z-buffer)
- ④ Accumulation buffer (images accumulated using a set of ROP)
  - e.g. motion blur, depth of field, antialiasing, soft shadows, ...
- ④ Double buffering: front buffer & back buffer, swapped during vertical retrace (avoid seeing uncompleted screen)

Digital ART Lab, SJTU



## Various Pipelines

- ④ Real-time rendering pipelines: decades of API and GPU evolution for real-time rendering applications
  - fixed-function pipeline (e.g. *Nintendo's Wii*, maybe the last)
    - On-off configuration
  - Programmable GPUs (the modern way!)
    - Program exactly operations in substages
- ④ Offline rendering pipelines: different evolution paths
  - Film rendering: commonly micropolygon pipelines
  - Academic, and predictive rendering applications (Pre-Viz): ray tracing renderers

Digital ART Lab, SJTU



## EVOLUTION OF GPU HARDWARE

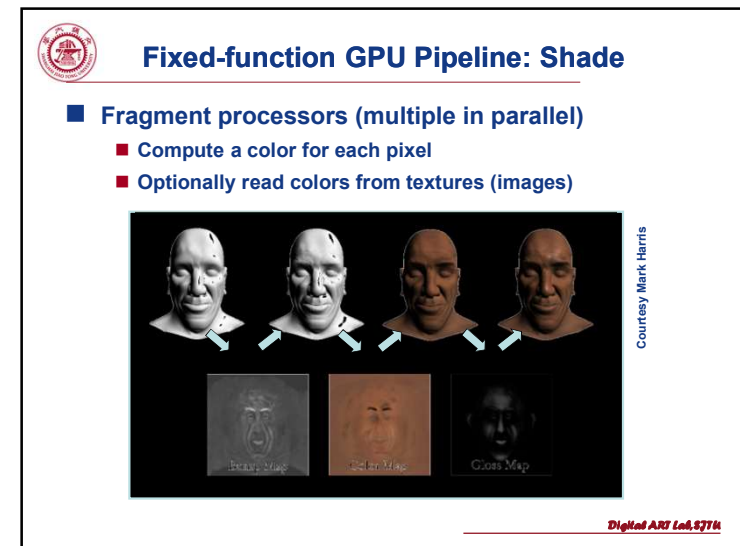
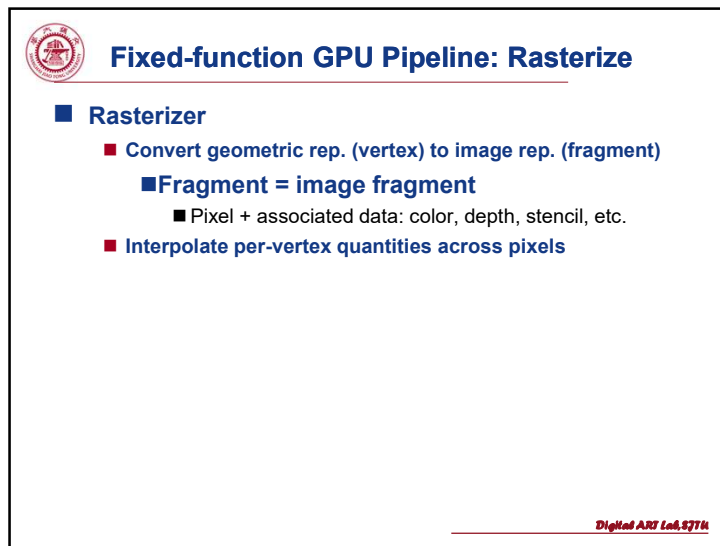
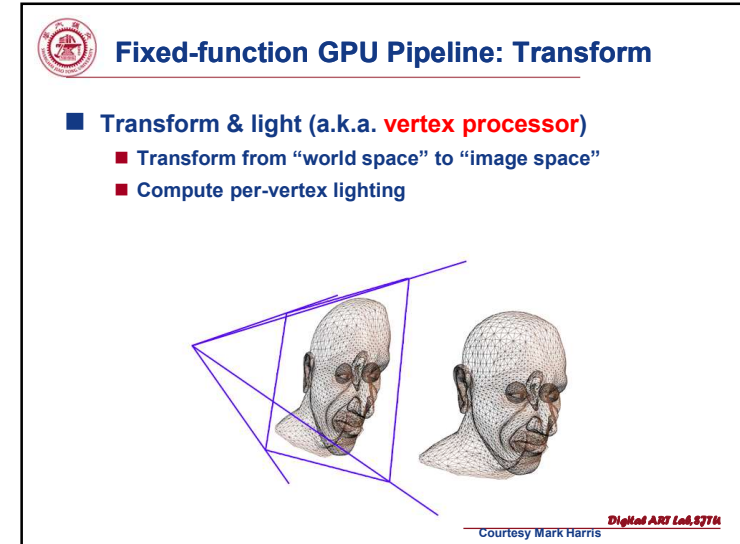
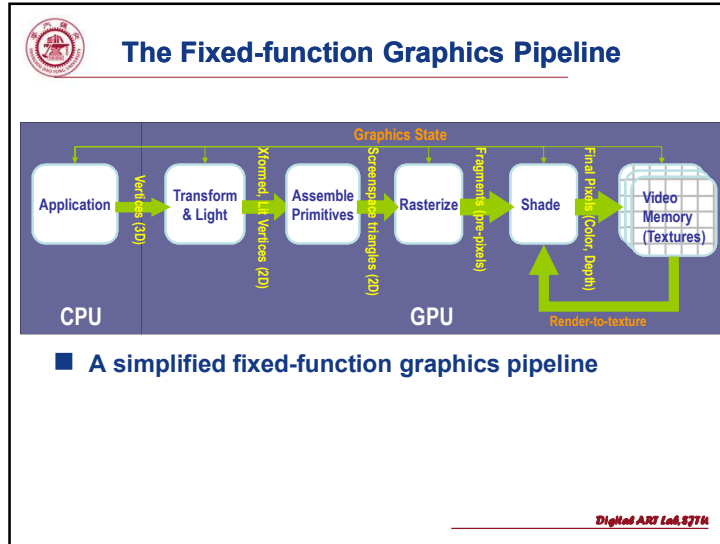
Digital ART Lab, SJTU

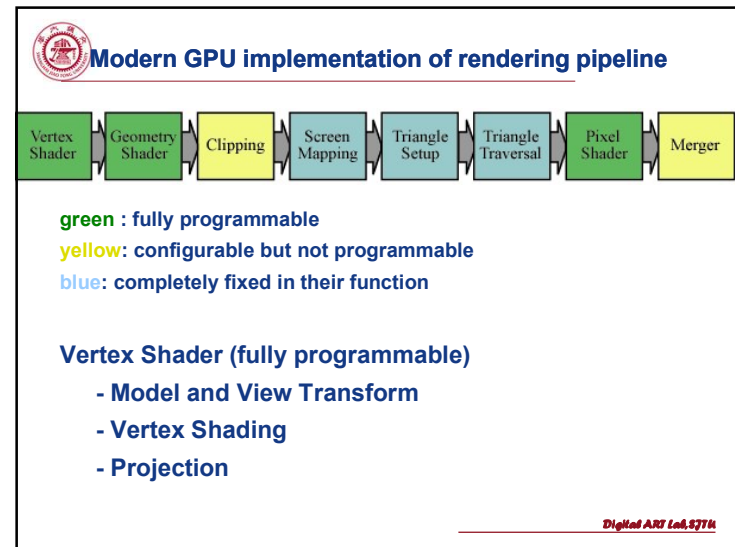
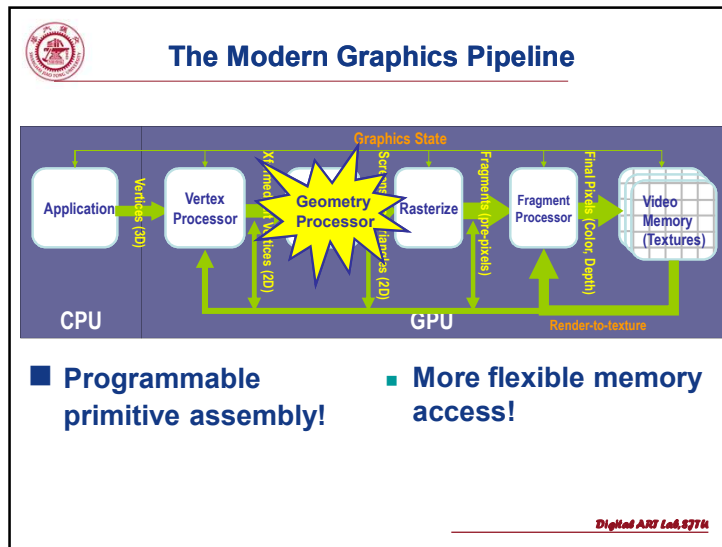
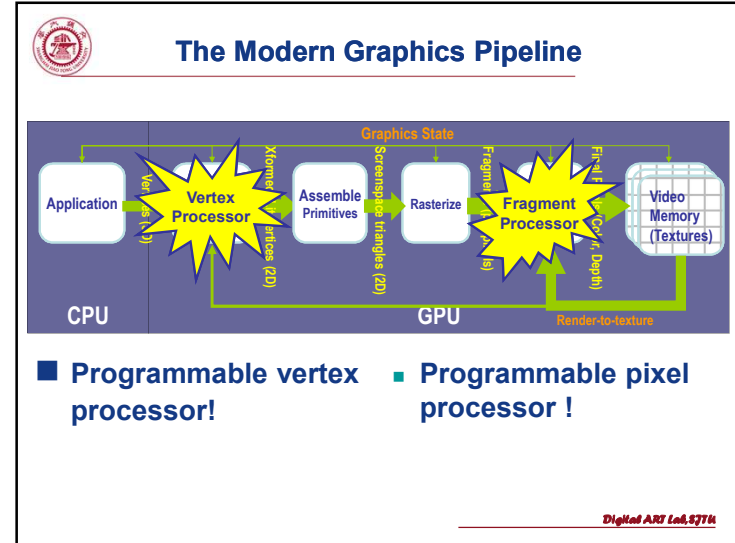
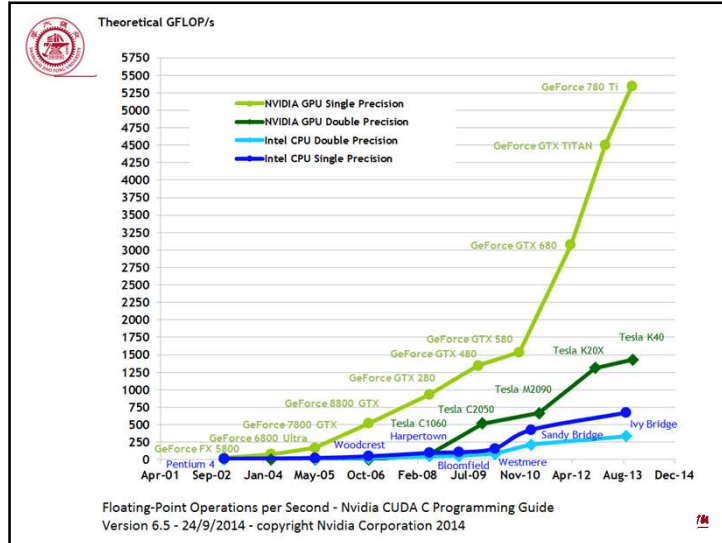


## Graphics Processing Unit (GPU)

- ④ Evolution of hardware graphics accelerations:
  - Started at the end of the pipeline
  - Worked back up the pipeline
  - Hardware accelerator for higher-level application-stage algorithms
- ④ 1999, NVIDIA GeForce256, coined the term GPU (vs. AMD VPU)
  - Hardware T&L (CPU vs. GPU) → **3D FPS Game Card** (bad for 2D apps)
  - 4 pixel pipelines: each has 4 pixel units + 1 texture unit
  - Triangle throughput: 15 Million/sec
  - Pixel throughput: 480 Million/sec
  - 2300 transistors ( > Pentium III)
  - 0.22 micro process (heat problem vs. 0.18 )
  - 256-bit display architecture
  - GeForce256 vs. Quadro

Digital ART Lab, SJTU





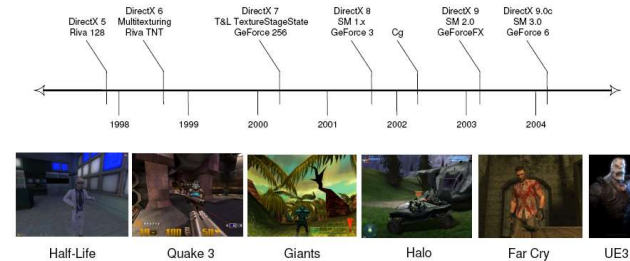


- ④ **Geometry shader (optional) - (fully programmable)**
  - Operate on the vertices of a primitive (point, line, or triangle)
- ④ **Clipping, screen mapping, triangle setup, triangle traversal stages (fixed)**
- ④ **Pixel Shading - (fully programmable)**
  - Pixel shading function
- ④ **Merge Stage (highly configurable)**
  - Modifying the color
  - Z-buffer blend, stencil, and other buffers

Digital ART Lab, SJTU



## DirectX Shader Model Timeline



Digital ART Lab, SJTU



## Earlier GPU Architecture Evolution

Product	New Features	OpenGL Version	Direct3D Version
2000 GeForce 256	Hardware transform & lighting, configurable fixed-point shading, cube maps, texture compression, anisotropic texture filtering	1.3	DX7
2001 GeForce3	Programmable vertex transformation, 4 texture units, dependent textures, 3D textures, shadow maps, multisampling, occlusion queries	1.4	DX8
2002 GeForce4 Ti 4600	Early Z culling, dual-monitor	1.4	DX8.1
2003 GeForce FX	Vertex program branching, floating-point fragment programs, 16 texture units, limited floating-point textures, color and depth compression	1.5	DX9
2004 GeForce 6800 Ultra	Vertex textures, structured fragment branching, non-power-of-two textures, generalized floating-point textures, floating-point texture filtering and blending	2.0	DX9c
2005 GeForce 7800 GTX	Transparency antialiasing	2.0	DX9c

Digital ART Lab, SJTU



## DirectX 5 / OpenGL 1.0 and Before

- **Hardwired pipeline**
  - Simple API inputs
  - Small set of operations
- **Example Hardware**
  - NVIDIA RIVA 128
  - 3dfx Voodoo
  - S3 Virge
- **Rigid data flow**
  - No read-back from frame buffer

Digital ART Lab, SJTU





### DirectX 6 / OpenGL 1.2

- Released 1998
- New Features
  - Multitexturing (in OpenGL since 1.3)
- Example Hardware
  - NVIDIA RIVA TNT
  - ATI Rage 128

Digital ART Lab, SJTU



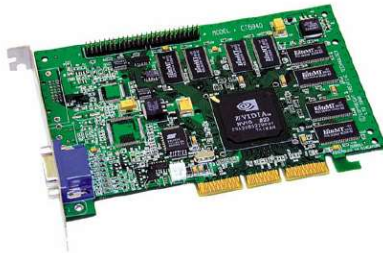
### DirectX 7 / OpenGL 1.3

- Released 1999
- More work outsourced to GPU
  - Hardware Transformation and Lighting (T&L, Direct3D only)
- New Features
  - Texture Compression in OpenGL
- Example Hardware
  - NVIDIA GeForce 256
  - ATI Radeon 7200

Digital ART Lab, SJTU



### Example: GeForce 256 (1999)



- Graphics Core:
  - 256-bit
- Memory Interface:
  - 128-bit
- Triangles/s:
  - 15 Million
- Pixels/s:
  - Up to 480 Million
- Memory:
  - Up to 128MB

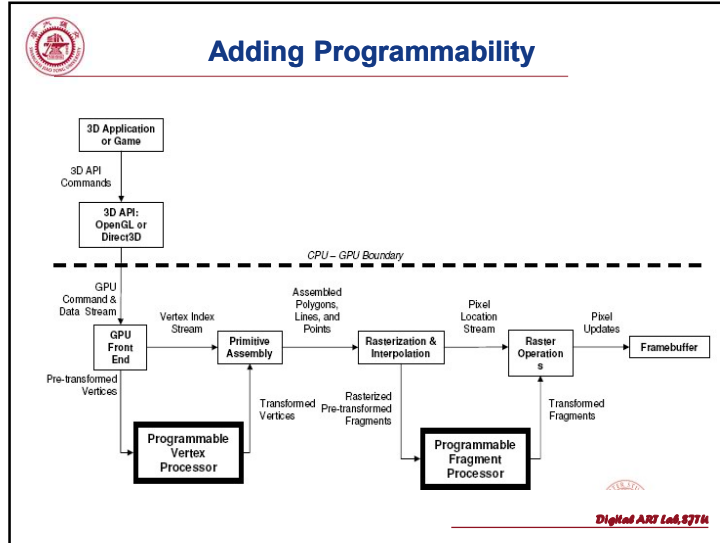
Digital ART Lab, SJTU



### DirectX 8

- Released 2000
- Introduction of Shader Model 1.1
  - Shaders are GPU-run programs that manipulate Vertices or Pixels
  - Enables a plethora of new visual effects
  - Adds programmable processors to the graphics pipeline
- Example Hardware:
  - NVIDIA GeForce 3
  - ATI Radeon 9000

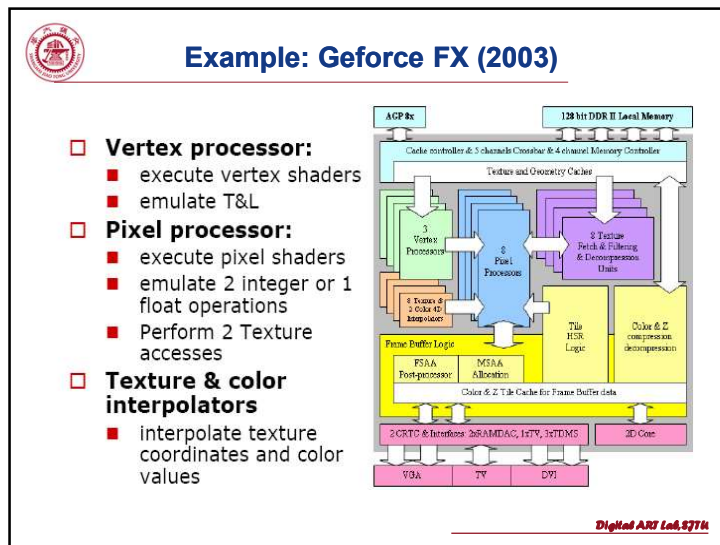
Digital ART Lab, SJTU



### DirectX 9

- Released 2002
- Much more general programming paradigm
  - Branching
  - Floating point fragment programming
- Shader Model 3.0 (in 9.0c, 2004)
  - Big feature increment
- Example Hardware
  - NVIDIA GeForceFX (9.0)
  - NVIDIA GeForce 6200 (9.0c)


*Digital ART Lab, SJTU*



### OpenGL 2.0

- Released 2005
- OpenGL Shading Language (GLSL)
  - Vertex and fragment shaders
  - GLSL ties shaders to OpenGL API
- Point sprites
  - Particle effects
- Many more features

*Digital ART Lab, SJTU*




## DirectX 10

---

- Released 2006
- Aligned with Windows Vista
- New Features
  - Geometry shaders; Streaming output; Arrays of surfaces and resource views; State encapsulation
- Break with Past
  - User mode drivers, even for DX9
  - Drivers do not implement Shader Models <4.0
  - No more fixed function (compiles legacy API calls to shaders)

Digital ART Lab, SJTU

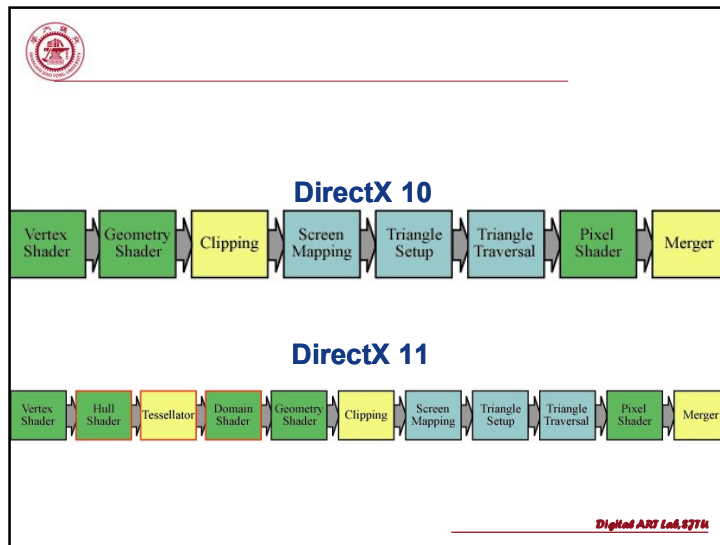



## DirectX 11

---

- Direct3d 11, Released 2009
  - Windows Vista(With Patch)/Windows 7
  - Shader Model 5.0
  - Tessellation, Multithreaded rendering, Compute shaders, supported by hardware and software running Direct3D 9/10/10.1
- ⊕ Direct3D 11.1
  - Windows 8, Stereoscopic 3D Rendering, GPGPU
- ⊕ Direct3D 11.2
  - Windows 8.1, Tiled resources, GPGPU

Digital ART Lab, SJTU



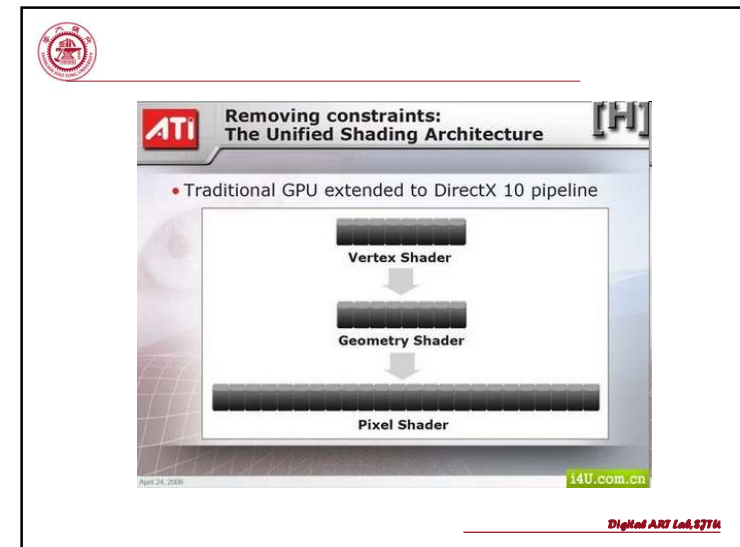
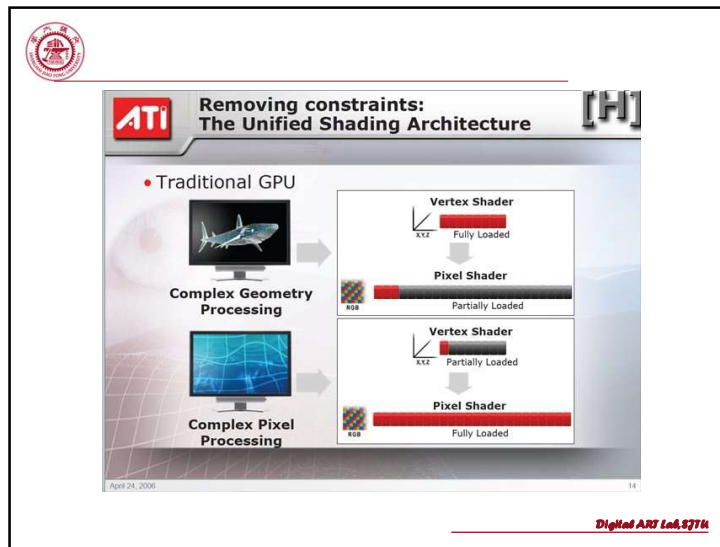
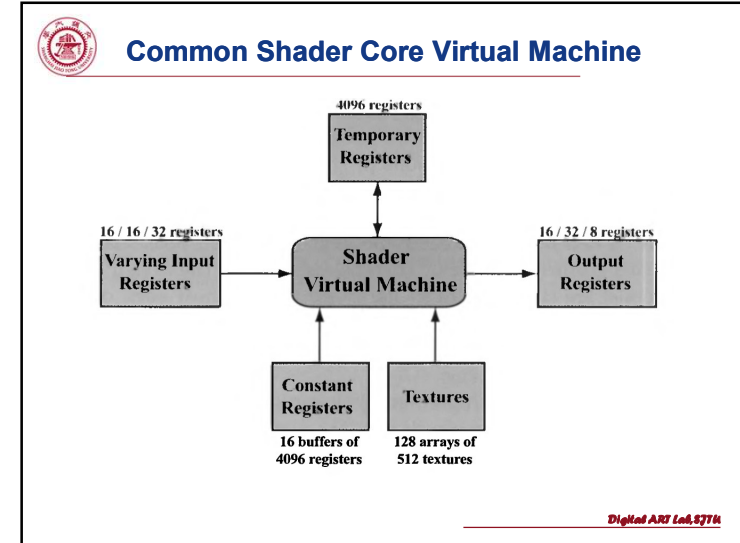
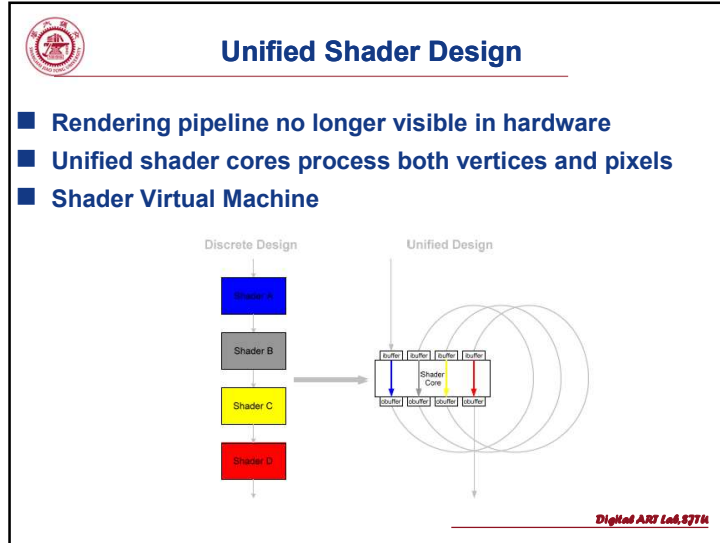


## DirectX 12

---

- ⊕ July 2015
- ⊕ Windows 10, Xbox One
- ⊕ reduce driver overhead: “console-level efficiency”
- ⊕ a lower level of hardware abstraction
  - enabling future games to significantly improve multithreaded scaling and (decrease) CPU utilization
- ⊕ claimed to be better than DirectX 11:
  - 50-70% faster
  - >50% reduction in power consumption

Digital ART Lab, SJTU





## Impact of Unified Shaders

- ❑ All shading processes performed by a unified set of processors
- ❑ Fewer bottle-necks (i.e. in case of vertex or pixel dominant scenes)
- ❑ Better hardware utilization
- ❑ Hardware architecture no longer reflects the graphics pipeline
- ❑ Greater flexibility makes GPUs eligible for non-graphics applications (game physics, scientific applications)

➔ Basically makes the GPU a massively parallel stream multiprocessor!

Digital ART Lab, SJTU



## Example of GPU Architecture



Digital ART Lab, SJTU



## Programmable Shader Stage

- ④ **Common-shader core (API) (after DX10)**
  - Vertex, pixel and geometry shaders share a programming model
  - Functional description seen by the application programmer
- ④ **Unified shaders (GPU architecture)**
  - A GPU architecture that maps well to the common-shader core
- ④ **Programming model:**
  - Shaders are programmed using C-like shading languages (e.g. Cg, HLSL, GLSL)

Digital ART Lab, SJTU



## ShaderToy Snail



- ④ GLSL pixel shader (~800 line codes)
- ④ Procedural modeling and procedural lighting
- ④ Raymarching
- ④ distancefield

```

188 return d/sc;
189
190 vec2 mapSnail( vec3 p, out vec4 matInfo )
191 {
192     vec3 head = vec3(-0.76,0.6,-0.3);
193     vec3 q = p - head;
194     // body
195     #if 1
196     vec4 b1 = sdBezier( vec3(-0.13,-0.65,0.0), vec3(0.24,0.94,0.1), head,vec3(0.0)
197     float d1 = b1.x;
198     d1 = smoothstep(0.0,0.2,b1.y)*0.16 - 0.07*smoothstep(0.5,1.0,b1.y));
199     b1 = sdBezier( vec3(-0.05,0.0,0.0), vec3(-0.3,0.3+0.05,0.0), head,vec3(0.06,-0.06,0.0)
200     float d2 = b1.x;
201     d2 = 0.1 - 0.04*b1.y;
202     d1 = min( d1, d2, 0.03 );
203     matInfo.xyz = b1.yzw;
204     #else
205     vec4 b1 = sdBezier( vec3(-0.13,-0.65,0.0), vec3(0.24,0.94,0.1), head,vec3(0.0)
206     float d1 = b1.x;
207     d1 = smoothstep(0.0,0.2,b1.y)*0.16 - 0.75*0.07*smoothstep(0.5,1.0,b1.y));
208     matInfo.xyz = b1.yzw;
209     float d2;
210     #endif
211     d2 = sdSphere( q, vec4(0.0,-0.06,0.0,0.05) );
212     d1 = min( d1, d2, 0.03 );
213     d1 = min( d1, sdSphere(p,vec4(0.05,0.52,0.0,0.13)), 0.07 );
214     q.xz = mat2(0.8,0.6,-0.6,0.8)*q.xz;
215     vec3 sq = vec3( q.xy, abs(q.z) );
216     // top antennas
217     vec3 a1 = 0.85*sin(0.5*time+vec3(0.0,1.0,3.0) + vec3(2.0,1.0,0.0)*sign(q.z) );
218     vec4 b2 = sdBezier( vec3(0.0), vec3(-0.1,0.2,0.2), vec3(-0.3,0.2,0.3)*sq );
219     float d3 = b2.x;
220
  
```

Digital ART Lab, SJTU